

NCIE-005 v1.2 → v1.3

Surgical Amendment Change Register

VPF & Dynamic Agent Synthesis Behavioral Consistency Surgical Amendment — IN DEVELOPMENT / FOR HUMAN REVIEW

Authorities: Approved NCIE-001 v1.3, approved NCIE-002 v0.5, approved NCIE-003 v0.4, approved NCIE-004 v1.2, and approved NCIE-007 v1.1. NCIE-006 v1.1 is used as the current Memory/Context baseline only (not surgically amended). The existing 35-chapter structure is unchanged; no chapter was added, removed or renumbered.

Amendment ID	Chapter/Section	Previous Behavioral Position (v1.2)	Surgical Amendment (v1.3)	Authority
AM-1	Ch.1, new §1.4	No statement of the current engineering/orchestration/canonical-data/Memory authority chain or explicit behavioral-ownership boundary.	Added §1.4 stating current authorities (NCIE-004 v1.2, NCIE-007 v1.1, NCIE-003 v0.4, NCIE-006 v1.1) and the boundary: NCIE-005 owns what ARGUS says/shows/asks; NCIE-004 owns physical implementation; NCIE-007 owns orchestration doctrine. Extended HR-1-1 to include NCIE-007 scoping.	NCIE-001 v1.3 NCIE-002 v0.5 NCIE-003 v0.4 NCIE-004 v1.2 NCIE-007 v1.1
AM-2	Ch.2, 2 new invariant-register rows, new §2.3	The ARGUS May / Does Not Autonomously register (Table 2) had no Agent Factory/Promotion rows; no Dynamic Agent/VPF invariant statements existed.	Added 2 rows to Table 2 (Dynamic Synthesis request / Promotion recommendation vs self-registration / Gate PASS-as-approval). Added §2.3 with 10 new invariants (AGENT FACTORY AVAILABLE ≠ SYNTHESIS AUTHORIZED, GATE A/B PASS ≠ HUMAN APPROVAL, VPF PASS ≠ FINDING/DECISION, MULTI-AGENT CONSENSUS ≠ HUMAN CONSENSUS, VALIDATED AI RESULT ≠ SOURCE EVIDENCE, VPF PASS ≠ SPOKEN-DISCLOSURE AUTHORIZATION, etc.), enforced by the same Tool/Model Gateway and VPF Runtime as the existing register.	NCIE-003 v0.4 §3.9; NCIE-004 v1.1 Ch.18-20; NCIE-007 v1.1
AM-3	Ch.7 §7.3; Ch.8 §8.3; Ch.10 §10.2; Ch.11 §11.3	Hypothesis, Contradiction and Evidence Gap behavior did not address AI/Agent origin or VPF state; the Output/Epistemic-Type chapter did not require VPF State/Qualification/AI Origin retention.	Ch.7: AI-generated hypotheses remain hypotheses regardless of origin/VPF state. Ch.8: material contradictions preserved through VPF/synthesis, never smoothed for PASS. Ch.10: Evidence Gap identification ≠ institutional confirmation; VPF validates the artifact, not the gap. Ch.11: substantive AI output retains applicable VPF State/Qualification/AI Origin/Human Modification State per the existing citation discipline.	NCIE-007 v1.1 (Ch.7,8,10); NCIE-003 v0.4 / NCIE-004 v1.2 / NCIE-006 v1.1 (Ch.11).
AM-4	Ch.13 §13.4; Ch.14 §14.3	Room/collaborative behavior did not address Agent participation, Agent-Room lifecycle, or Agent convergence vs Room/human consensus.	Ch.13: Agents may contribute to Rooms without becoming members/officers; provenance retained; Agent convergence ≠ Room consensus; Ephemeral Agents don't persist in a Room by default. Ch.14: Human Contribution/ARGUS/Agent/Dissent/Consensus/Human Judgment remain distinguishable; Agent dissent is an AI-originated challenge, not automatically a Finding; Human Judgment stays human.	NCIE-007 v1.1
AM-5	Ch.15 §15.3, new HR-15-3	Evidence-Watch re-entry triggers covered Evidence-Watch events only; no VPF revalidation, Agent dependency-impact or Agent restriction/suspension trigger existed.	Added §15.3 extending re-entry triggers to VPF Revalidation, Agent Dependency Impact and Registered Agent Restriction/Suspension, applying the existing §15.2 materiality/anti-fatigue discipline. Added HR-15-3 (notification threshold tuning for the new triggers).	NCIE-003 v0.4 §3.9; NCIE-004 v1.2
AM-6	Ch.21 §21.4	High-consequence Anti-Fraud/SIMBOX assistance did not explicitly exclude Agent Confidence/Consensus/Gate PASS as substitutes for required Human/Institutional authority.	Added §21.4 stating Agent Confidence, Agent Consensus, Gate A PASS, Gate B PASS and Prior Agent Success are not substitutes for the Human/Institutional authority this chapter already requires for tracking/blocking/Decision-adjacent action.	NCIE-007 v1.1
AM-7	Ch.23 §23.2	Memory/Decision-Time behavior did not distinguish Task-Scoped Agent Context from Institutional Memory/Conversational Context, nor state MEMORY ≠ CURRENT TRUTH for Agent/Tool/Model eligibility.	Added §23.2: four context types kept distinct; Agent Context is minimum-necessary and its absence is not fabricated; task-scoped Context expires with the Run (AGENT RUNTIME STATE ≠ INSTITUTIONAL MEMORY); MEMORY ≠ CURRENT TRUTH for consequential current-state questions; Decision-Time answers use historical-cutoff state only.	NCIE-006 v1.1 NCIE-007 v1.1

Amendment ID	Chapter/Section	Previous Behavioral Position (v1.2)	Surgical Amendment (v1.3)	Authority
AM-8	Ch.24 §24.2	Institutional Knowledge behavior did not address Validated AI Result vs Knowledge promotion, or retracted-Knowledge dependency impact.	Added §24.2: VALIDATED AI RESULT ≠ INSTITUTIONAL KNOWLEDGE; VPF PASS ≠ KNOWLEDGE PROMOTION; AI synthesis never presented as the underlying Knowledge/Evidence source; retracted/superseded Knowledge triggers the applicable impact assessment before dependent AI output continues being presented as current.	NCIE-003 v0.4 Ch.22; NCIE-0 v1.1.
AM-9	Ch.25 (principal amendment), new §25.2-§25.6, new HR-25-1, HR-25-2	Tool/action UX covered Tool proposal/execution/approval only; no Capability Discovery, Registered-Agent-reuse, Dynamic Synthesis, Agent validation/execution status, self-registration/recursive-synthesis refusal, or Promotion/Engineering-Change-Candidate UX existed.	Added §25.2 (scope extension), §25.3 (capability discovery + synthesis status language, authoritative-state-only status), §25.4 (authority/self-registration-prohibited/recursive-synthesis-disabled/Codex boundary), §25.5 (Tool proposal ≠ execution, approval-state distinctions, OUTCOME UNKNOWN, no blind retry), §25.6 (Gate A UX, failure handling, technical-detail minimization). Added HR-25-1 (status-stage terminology) and HR-25-2 (Promotion Candidate approval ownership — Blocking).	NCIE-001 v1.3 §12.77-§12.80; NCIE-003 v0.4 §3.9; NCIE-00 Ch.20; NCIE-0 v1.1.
AM-10	Ch.26 §26.2	Output/deliverable behavior did not require retention of AI Origin/Agent Origin/VPF State/Qualification, nor address composite-output VPF inheritance.	Added §26.2: outputs preserve applicable AI Origin, Agent Origin, VPF State, Qualification, Evidence References and Human Modification State; VPF-PASS ≠ Institutionally Approved/Official Finding/Human-Verified; QUALIFIED and Human-Review-required outputs remain visibly distinguishable; composite reports don't automatically inherit component PASS (COMPONENT PASS ≠ COMPOSITE PASS).	NCIE-007 v1.1
AM-11	Ch.27 §27.2	Provenance/attribution behavior predated Agent/VPF origin distinctions and human-modification-of-AI-artifact provenance.	Added §27.2: provenance distinguishes ARGUS Direct AI, Registered Agent, Ephemeral Agent, Multi-Agent Composite, VPF Validation, Human Modification and Human Decision/Approval, at a level supporting trust/review/Decision-Time reconstruction without raw system traces; Evidence citation discipline is not weakened by Agent origin; human-edited AI artifacts retain the Original/Modified distinction.	NCIE-003 v0.4 NCIE-004 v1.2 NCIE-007 v1.1
AM-12	Ch.28 §28.2	Personalization behavior did not state a boundary against weakening Agent/VPF/Tool governance via user preference.	Added §28.2: personalization may change tone/detail/presentation/notification but not Authorization, VPF, Agent Risk Controls, Tool Approval, Evidence Requirements, Security or Human-Primary Authority; QUALIFIED/HUMAN REVIEW REQUIRED wording may not be softened into PASS/optional for convenience; no preference can override mandatory Agent/Tool/VPF/Human Approval controls.	NCIE-007 v1.1
AM-13	Ch.29 (principal amendment), new §29.2-§29.3, new HR-29-3	Voice/privacy/safety behavior addressed protected-identity disclosure but not VPF-PASS-as-spoken-authorization, Agent-sourced spoken content, or AI-content claiming its own VPF state (prompt-injection/VPF-manipulation).	Added §29.2: VPF PASS ≠ SPOKEN-DISCLOSURE AUTHORIZATION; voice disclosure controls apply identically regardless of AI/Agent source; Voice Presence/Speaker Recognition/Transcript/spoken 'yes' is not approval absent authenticated mechanism; spoken status preserves semantic accuracy (no 'Approved' for mere Gate/QUALIFIED state); Agent telemetry/context not exposed via voice. Added §29.3: instructions embedded in Evidence/Memory/External Content/Tool/Agent Output are treated as data unless from the governed instruction path; AI text claiming 'VPF PASS' is not authoritative. Added HR-29-3 (Agent-originated spoken-disclosure policy, Blocking).	NCIE-003 v0.4 §3.9; NCIE-00
AM-14	Ch.30 §30.2, new HR-30-1	Degraded-mode behavior had no Agent Factory/Ephemeral Runtime/VPF/Registry unavailability states, and no Human Review Register item at all for this chapter.	Added §30.2 distinguishing Agent Factory/Ephemeral Runtime/VPF/No-Eligible-Model/Tool/Registry/Human-Review unavailability from each other and from total ARGUS failure; VPF unavailable never fails open; AI/AGENT/VPF FAILURE ≠ TOTAL NCIE FAILURE. Added HR-30-1 (degraded-mode notification threshold policy, Blocking — first Human Review item for this chapter).	NCIE-004 v1.2 Ch.30; NCIE-0 v1.1.
AM-15	Ch.31 §31.2	Behavioral evaluation did not test Dynamic Agent/VPF scenarios (synthesis, reuse, capability-gap, self-registration/recursive-synthesis	Added §31.2 specifying the full Dynamic Agent/VPF behavioral test suite (13 scenario classes) mirrored from directive §123-142; tests are not marked passed without implementation evidence.	NCIE-007 v1.1

Amendment ID	Chapter/Section	Previous Behavioral Position (v1.2)	Surgical Amendment (v1.3)	Authority
		refusal, VPF bypass/manipulation, degraded modes, Decision-Time).		
AM-16	Ch.32 §32.2	NFR chapter did not address Agent/VPF-specific responsiveness, graceful degradation, or accessibility for new status indicators.	Added §32.2: ARGUS stays responsive with truthful state during Agent/VPF processing (no invented latency targets); unaffected capability preserved when components fail; new status/qualification indicators meet existing accessibility requirements; authoritative state is not described inconsistently across channels.	NCIE-004 v1.2
AM-17	Ch.33 §33.2	Configuration governance did not address Agent/VPF-facing settings or state that configuration cannot grant authorization.	Added §33.2: configuration extends to Dynamic Agent UX/status/VPF presentation settings; ARGUS CONFIGURATION ≠ INSTITUTIONAL AUTHORIZATION — no preference may enable VPF bypass, self-registration, recursive synthesis, or unauthorized Tools/data/disclosure; presentation-only configuration cannot alter the authoritative VPF Profile/Disposition or Agent Authorization/Tools/Models/Context/Risk/Runtime Limits.	NCIE-004 v1.2 Ch.33; NCIE-0 v1.1.
AM-18	Ch.34 (mechanical reconciliation)	Human Review Register totaled 17 items (8 Proposed / 6 Institutional / 2 Legal / 1 Resolved / 0 Workflow; 8 Blocking / 9 Non-Blocking).	Mechanically re-extracted after all substantive amendments: 22 items (10 Proposed / 8 Institutional / 3 Legal / 1 Resolved / 0 Workflow; 11 Blocking / 11 Non-Blocking). 5 new items added (HR-15-3, HR-25-1, HR-25-2, HR-29-3, HR-30-1), inserted in correct chapter order in Table 31; 0 removed. Table 30 summary recomputed programmatically from final Table 31 content. Added §34.6 Reconciliation Note recording the delta arithmetic (17 + 5 - 0 = 22).	Administrative directive §156- (Chapter 34 Human Review reconciliation).
AM-19	Ch.35, new §35.1A, §35.2A, §35.5, 23 traceability rows updated	Traceability matrix cited no NCIE-007 authority at all (predates NCIE-007 as a separate document) and did not address Dynamic Agent/VPF acceptance or a behavioral-readiness disposition.	Added §35.1A (current authority chain, explicitly including NCIE-007 v1.1 and NCIE-006 v1.1 as current baseline, not falsely amended). Added §35.2A (Dynamic Agent/VPF acceptance: ARGUS-requests/Agent-Factory-synthesizes/Runtime-executes relationship; Gate A/B never Human Approval; no VPF-bypass route; self-registration/recursive-synthesis refused; BEHAVIOR SPECIFIED ≠ CAPABILITY IMPLEMENTED). Added §35.5 (Readiness Disposition: Generated-Code NOT presented as available; Recursive-Synthesis disabled by default; Self-Registration prohibited; explicit downstream handover to NCIE-006/007/008/009/011/012/016). Updated 23 traceability-matrix rows (Ch.1,2,7,8,10,11,13-15,21,23-30 and 32-35) with their NCIE-003/004/006/007 dependency.	Approved NCIE v1.3 / NCIE-006 v1.1 / NCIE-003 v0.4 / NCIE-004 v1.2 / NCIE-007 v1.1 / NCIE-006 v1.1 current baseline
AM-DOC	Cover / Document Governance / Running Header	Cover cited NCIE-002 v0.4/NCIE-003 v0.3/NCIE-004 v1.1 with no NCIE-007/006 mention. Document Governance table's Edition field read 'In development — Version 1.0' — the same recurring pre-existing template defect already found and fixed in NCIE-004 v1.2's own amendment (cover already said 'Version 1.2'). The running header read 'VERSION 1.2' on all 54 rendered pages after every substantive amendment had been applied.	Cover corrected to cite approved NCIE-001 v1.3 / NCIE-002 v0.5 / NCIE-003 v0.4 / NCIE-004 v1.2 / NCIE-007 v1.1, with NCIE-006 v1.1 identified as the current Memory/Context baseline (not an amendment authority). Governance table's Edition corrected to 'Version 1.3'; Requirements/Architecture/Canonical-data/Engineering-baseline rows corrected to v1.3/v0.5/v0.4/v1.2 respectively; new Orchestration authority (NCIE-007 v1.1), Memory/Context baseline (NCIE-006 v1.1) and Prior baseline (v1.2) rows inserted. New §0.10 Document Amendment History added. Running header corrected to 'VERSION 1.3' and verified across all 54 rendered pages before delivery, checked proactively this time based on the identical defect found in the two prior NCIE-004 amendment passes.	Administrative directive §436-§493 (version/document control hard gate §457 (page-level verification evidence), the specific gates the directive exists to enforce.

Chapters Reviewed, No Substantive Change Required

Chapters 3, 4, 5, 6, 9, 12, 16, 17, 18, 19, 20, 22 were reviewed against the current approved authority chain per directive §41-42, §56 (Chapters 3-4 preservation; Chapters 5-12 analytical behavior; Chapters 16-22 domain assistance). No direct contradiction with the current approved architecture was found, and no genuinely new behavioral consequence was material enough to warrant amendment — consistent with the directive's instruction not to manufacture behavior merely to touch every chapter.

Human Review Register Reconciliation

17 (v1.2) + 5 added – 0 removed = 22 (v1.3). Category deltas: Proposed Design Default 8→10 (+2: HR-15-3, HR-25-1); Institutional Confirmation Required 6→8 (+2: HR-25-2, HR-30-1); Legal/Policy Confirmation Required 2→3 (+1: HR-29-3); Resolved by Upstream Baseline unchanged at 1; Workflow/Source Discovery Required unchanged at 0. Blocking 8→11 (+3: HR-25-2, HR-29-3, HR-30-1); Non-Blocking 9→11 (+2: HR-15-3, HR-25-1).

Additional Consistency Amendments Found (Full-Document Contradiction Scan)

A full-document scan for stale current-authority version references (NCIE-001 v1.2, NCIE-002 v0.4, NCIE-003 v0.3, NCIE-004 v1.1 cited as current; NCIE-005 v1.2 cited as the current edition) found two occurrences requiring correction: the Document Governance table's pre-existing 'Edition: Version 1.0' defect and the running header's 'VERSION 1.2' on all 54 pages (both AM-DOC). No other pre-existing stale or contradicting statement was found. A semantic scan for statements implying Agent Synthesis = Authority, Agent Success = Approval, VPF PASS = Human Approval/Finding/Decision, Agent Consensus = Human Consensus, self-registration availability, or recursive-synthesis availability found none; every governing invariant from the directive's §9 list is stated explicitly and preserved throughout. No additional surgical amendment beyond AM-1 through AM-19 and AM-DOC was required.