

NCIE-006 v1.1 → v1.2

Surgical Amendment Change Register

VPF & Dynamic Agent Context Surgical Consistency Amendment — IN DEVELOPMENT / FOR HUMAN REVIEW

Authorities: Approved NCIE-001 v1.3, approved NCIE-002 v0.5, approved NCIE-003 v0.4, approved NCIE-004 v1.2, approved NCIE-005 v1.3, and approved NCIE-007 v1.1. The existing 36-chapter structure is unchanged; no chapter was added, removed or renumbered.

Amendment ID	Chapter/Section	Previous Memory/Context Position (v1.1)	Surgical Amendment (v1.2)
AM-1	Ch.1, new §1.4	Authority chain cited NCIE-001 v1.2/NCIE-002 v0.4/NCIE-003 v0.3/NCIE-004 v1.1/NCIE-005 v1.2 as current; no explicit Agent Context/Decision-Time/dependency boundary statement.	Added §1.4 stating the current authority chain (NCIE-004 v1.2, NCIE-005 v1.3, NCIE-007 v1.1) and explicitly disclaiming the older versions as current. Restates NCIE-006's boundary: what Context an Agent may receive, how bounded, provenance references, expiry, Decision-Time participation, dependency/triggers — not Agent Factory/runtime implementation.
AM-2	Ch.2, 3 new terms in Table 2, new §2.4	12 Governing Invariants (front matter §0.3) with no Agent/VPF-specific extensions; no 'Agent Context' terms in the authoritative definitions table.	Added Agent Context, Agent Task Context and Agent Run Context to Table 2's definitions. Added §2.4: 10 new Agent/VPF invariants extending the existing epistemic boundaries (AGENT RUNTIME STATE, INSTITUTIONAL MEMORY, AGENT CONTEXT, AGENT AUTHORITY, EPHEMERAL AGENT EXPIRY ≠ REQUIRED PROVENANCE DELETION, VALIDATED AI RESULT ≠ INSTITUTIONAL KNOWLEDGE, VPF PASS ≠ KNOWLEDGE PROMOTION, HISTORICAL AGENT/VPF STATE, CURRENT ELIGIBILITY, RECOVERY ≠ REAUTHORIZATION, SEMANTIC RETRIEVAL, AUTHORITY TO DISCLOSE, AI ANALYSIS OF EVIDENCE ≠ EVIDENCE ITSELF, VPF VALIDATION STATE ≠ HUMAN APPROVAL STATE). Implements the mandatory NCIE-003 v0.4 handover for Agent Task Context, Agent Run Context, VPF Artifact References, Promotion History and Context Expiry.
AM-3	Ch.3 §3.3	Memory class taxonomy did not address whether Dynamic Agent execution creates a new Memory class.	Added §3.3: Agent execution does not create a permanent 'Agent Memory' class; temporary state classified as Task/Run Context or Ephemeral Run State using the existing Memory-class lifecycle.
AM-4	Ch.4 §4.4	Context Envelope (§4.1) had no Agent Run Reference or Authorization Context field, and no explicit CONTEXT ENVELOPE ≠ AUTHORIZATION GRANT statement.	Added §4.4 extending the Envelope with Agent Run Reference and Authorization Context; states CONTEXT ENVELOPE ≠ AUTHORIZATION GRANT, MINIMUM NECESSARY AGENT CONTEXT, purpose binding, and expiry that preserves required Audit/Provenance/Run History/VPF History/Legal material.
AM-5	Ch.5 §5.3 (covering Ch.5-10 consolidated review)	Conversation/Preference/Workspace/Room/Investigation/Situation Memory chapters had no Agent-participation boundary statement.	Added §5.3 stating Agent participation does not create permanent Agent-owned Memory, with the chapters specific invariants for Ch.6-10: PREFERENCE ≠ PERMISSION; Workspace/Session isolation; ROOM MEMBERSHIP ≠ CROSS-ROOM ACCESS; Investigation-bounded Agent Context; Situation context domain synthesis ≠ cross-domain permission. Chapters 6-10 themselves reviewed with no additional chapter specific change needed beyond this consolidated statement.
AM-6	Ch.11 §11.3 (covering Ch.11-13)	Decision Memory/Decision-Time/Current-State chapters had no Agent Run/Result/VPF provenance reference or historical-eligibility-≠-current-eligibility statement.	Added §11.3: Decision Memory may reference Agent Run, Agent Result Package, VPF Validation Artifact and Human Decision as distinct objects (AI/AGENT/VPF HISTORY ≠ HUMAN DECISION AUTHORITY); Decision-Time reconstruction (Context uses the historical Agent Definition Version/Run/Context/Model Route/Tool Outcomes/Result/VPF State at the cutoff, never later).

Amendment ID	Chapter/Section	Previous Memory/Context Position (v1.1)	Surgical Amendment (v1.2)
			state; Current-State Reconciliation (Ch.13) resolve current Agent/Model/Tool/VPF eligibility from authoritative sources (HISTORICAL ≠ CURRENT ELIGIBILITY).
AM-7	Ch.14 §14.2 (principal amendment)	Context Assembly supported ARGUS retrieval generally; no explicit Agent Run/Ephemeral Agent/VPF Validation assembly path, inputs list, or fail-safe rule.	Added §14.2: assembly extended to ARGUS Direct Registered Agent Runs, Ephemeral Agent Runs and VPF Validation; inputs (Task Contract, Effective Principal, Purpose, Agent Definition, current Authorization, Memory/Evidence/Knowledge Score, Temporal Cutoff, Classification, Resource/Context Limits); output attributable to the exact Agent Run. CONTEXT SELECTION ≠ TASK AUTHORIZATION. AGENT DEFINITION / TASK CONTRACT ≠ CONTEXT AUTHORITY; fail-safe (no broader Context by default if authorization/scope unclear).
AM-8	Ch.15 §15.2	Context Propagation addressed service/workflow propagation generally; no Agent-to-Agent delegation Context bounding or child-Context-minimization rule.	Added §15.2: child Context bounded by the interests of effective-principal authorization, parent delegation scope, child task, child Agent Definition and current Context policy; parent's full Envelope is never copied automatically; lineage preserved; RECURSIVE DYNAMIC AGENT SYNTHESIS REMAINS DISABLED (Context propagation must not grant Factory authority).
AM-9	Ch.17 §17.3	ARGUS Memory Retrieval Interface addressed ARGUS retrieval; no explicit Agent-retrieval or AI-artifact-retrieval provenance rule.	Added §17.3: Agent retrieval bounded identically to ARGUS retrieval (Authorization/Purpose/Classification/Context/Temporal Scope/Provenance); RETRIEVABLE ≠ AUTHORITY TO DISCLOSE; retrieved AI artifacts preserve AI Origin/Agent Origin/VPF State/Qualification/Human Modification State and are never surfaced as source Evidence.
AM-10	Ch.18 §18.4	Institutional Knowledge Candidate Lifecycle had no Validated-AI-Result/Agent-Result/human-modification boundary against automatic Knowledge promotion.	Added §18.4: VALIDATED AI RESULT ≠ INSTITUTIONAL KNOWLEDGE; VPF PASS ≠ KNOWLEDGE PROMOTION; AGENT RESULT ≠ KNOWLEDGE CANDIDATE BY DEFAULT; Knowledge Candidate provenance (Source Evidence/Artifact/Agent Run/VPF Validation/Human Promotion) kept distinct; human modification before promotion preserves original vs modified provenance; Knowledge retraction triggers impact assessment for dependent output.
AM-11	Ch.22 §22.4	Dependency Graph & Downstream Impact Tracking had no Agent/VPF-specific dependency types or no-automatic-universal-invalidation rule.	Added §22.4: dependency graph extended to Evidence/Knowledge/Agent Definition/Model/Tool/Context/VPF Profile changes. DEPENDENCY ≠ AUTHORITY; no automatic universal invalidation — targeted impact analysis; historical Validation Artifacts preserved regardless of downstream changes.
AM-12	Ch.25 §25.4-§25.5	Memory Security/Zero-Trust chapter addressed human/service identity; no Agent Runtime Identity class, context-exfiltration boundary, or context-at-rest/cache/poisoning rule.	Added §25.4: Zero-Trust extended to Agent Runtime Identity (DELEGATED TASK ≠ HUMAN CREDENTIAL DELEGATION, AGENT IDENTITY ≠ AUTHORITY, current not cached authorization); CONTEXT EXFILTRATION BOUNDARY (reading/exporting/speaking/Tool-transmitting/re-disclosure) ≠ SYSTEM ADMINISTRATION ≠ CONTENT AUTHORITY. Added §25.5: runtime isolation, at-rest/in-transit protection, cache-reuse re-authorization, integrity/poisoning/contamination controls extend to Agent Context.
AM-13	Ch.26 §26.3-§26.4	Sensitive Memory/Protected Identity/Voice Context chapter had no explicit VOICE CONTEXT ≠ AUTHENTICATED IDENTITY statement extended to Agents, nor an Agent-specific protected-identity boundary.	Added §26.3: VOICE CONTEXT ≠ AUTHENTICATED IDENTITY; VPF PASS ≠ SPOKEN-DISCLOSURE AUTHORIZATION applicable to Agent-initiated retrieval identically to ARGUS retrieval. Added §26.4: protected-identity

Amendment ID	Chapter/Section	Previous Memory/Context Position (v1.1)	Surgical Amendment (v1.2)
			access/leakage/inference controls apply identically to Agent Context, with no broadening for Dynamic Agent/VPF capability.
AM-14	Ch.27 §27.3	Memory Provenance/Attribution/Audit chapter had no Agent/VPF-specific end-to-end provenance-reconstruction statement.	Added §27.3: an authorized reviewer can reconstruct Request → Context Assembly → Agent Run → Evidence/Knowledge Dependencies → Model/Toolchain Provenance → Agent Result → VPF Validation → ARGUS Presentation → Human Decision/Modification without private chain-of-thought; TELEMETRY ≠ AUDIT ≠ PROVENANCE preserved for Agent Context/VPF/human-modification audit specificities.
AM-15	Ch.28 §28.4	Retention/Legal Hold/Archival chapter had no Ephemeral-Agent-specific retention/disposition rule.	Added §28.4: EPHEMERAL AGENT ≠ EPHEMERAL PROVENANCE (expiry does not erase required Agent Definition/Run/Task Contract/Result/VPF/audit/Decision-Time records; runtime scratch-state cleanup distinct from Agent Context/Validation Artifact/Qualification/Promotion/Engineering-Change Candidate retention; DELETION ≠ HISTORICAL FALSIFICATION).
AM-16	Ch.29 §29.3	Memory Storage Architecture chapter had no explicit statement that Agent/VPF canonical references are logical, not separately-stored.	Added §29.3: LOGICAL ENTITY ≠ SEPARATE DATABASE for Agent Definition/Run/Task Contract/Result Package/VPF Profile/Validation Artifact references; reuses existing approved persistent architecture per NCIE-004 v1.2 Ch.33 §33.4's reusable Registry review; NCIE-006 references these objects; does not own their canonical storage.
AM-17	Ch.30 §30.3	Memory APIs/Events chapter had no Agent/VPF lifecycle event rule.	Added §30.3: material Agent/VPF lifecycle events extend the existing event architecture; EVENT ≠ CURRENT AUTHORITY STATE (events trigger reconciliation but do not replace current Registry/authorization/VPF sources or establish institutional authority).
AM-18	Ch.31 §31.3	Scalability/Performance/Observability chapter had no Agent Context telemetry-privacy rule.	Added §31.3: Agent Context telemetry observable extends existing architecture, subject to current authorization/purpose/classification/minimum-necessary-scope for telemetry privacy; no broader access via Dynamic Agent/VPF capability.
AM-19	Ch.32 §32.3	Resilience/Backup/Recovery chapter had no Agent Context/Agent-eligibility/VPF-state post-restore reconciliation rule.	Added §32.3: RECOVERY ≠ REAUTHORIZATION; post-restore reconciliation applies to Agent Context (expired stays expired), Agent eligibility (current Registry prevails), VPF state (obsolete Profiles do not reactivate) and vector/search indexes (current deletion/restriction/classification/authorization respected); deterministic continuity preserved where Agent/VPF integration is degraded.
AM-20	Ch.33 §33.3	Legacy Memory Migration chapter had no provider-native-memory or migration-integrity rule for Agent/VPF references.	Added §33.3: PROVIDER MEMORY ≠ NCIE INSTITUTIONAL MEMORY (provider-native conversation/Agent/session memory is implemented by capability unless explicitly governed and mapped); MIGRATED / IMPORTED ≠ APPROVED / ELIGIBLE; migration preserves identity/version/integrity/provenance/classification reports (does not fabricate) missing historical metadata.
AM-21	Ch.34 §34.3	Memory & Context Testing chapter had no Agent/VPF-specific test catalogue.	Added §34.3: minimum Agent/VPF Memory/Context test suite (Agent Context isolation, minimum-necessary assembly, current/remembered authorization, Agent expiry, Result-vs-Knowledge distinction, Evidence/Knowledge dependency-impact, Decision-Time historical-state-only reconstruction, current-reconciliation, semantic-similarity/deleted-embedded disclosure, prompt injection, voice/spoken-disclosure, human-modification-after-PASS, multi-Agent/child Agent isolation, cross-task cache authorization, recovery); TEST SPECIFIED ≠ TEST EXECUTED.

Amendment ID	Chapter/Section	Previous Memory/Context Position (v1.1)	Surgical Amendment (v1.2)
AM-22	Ch.35 (mechanical reconciliation)	Human Review Register totaled 16 items (1 Resolved / 2 Proposed / 6 Institutional / 2 Workflow / 4 Legal / 1 Security; 13 Blocking / 3 Non-Blocking).	Mechanically re-extracted: 20 items (1 Resolved / Proposed / 9 Institutional / 2 Workflow / 4 Legal / Security; 17 Blocking / 3 Non-Blocking). 4 new items added (HR6-4-2 Agent Context retention, HR6-18 Knowledge Candidate review steps, HR6-22-1 dependency-impact ownership, HR6-33-2 provide native memory governance), inserted in correct chapter order; 0 removed. Table 16 summary and Table 17 register recomputed programmatically from final content. Added §35.6 Reconciliation Note (16 + 4 = 20).
AM-23	Ch.36, new §36.5-§36.6, 23 traceability rows updated	Traceability matrix cited no NCIE-007 authority (predates NCIE-007 as a separate document) and did not address Dynamic Agent/VPF acceptance.	Added §36.5 (current authority chain including NCIE-007 v1.1, explicitly disclaiming v1.1's older citation as current). Added §36.6 (Dynamic Agent/VPF acceptance: NCIE-006 owns Memory/Context only; does not absorb Agent Factory/Registry/VPF/IAM/storage specification; bounded Context assembly/propagation; Decision integrity; MEMORY ≠ CURRENT TRUTH; Evidence/Knowledge boundaries; TEST SPECIFICATION TEST EXECUTED). Updated 23 traceability rows (Ch.1,2,3,4,5,11,14,15,17,18,22,25-33,34,35,36) with their NCIE-003/004/005/007 dependency.
AM-DOC	Cover / Document Governance / Running Header	Cover and Document Governance table cited NCIE-002 v0.4/NCIE-003 v0.3/NCIE-004 v1.1/NCIE-005 v1.2 with no NCIE-007 mention (Edition field itself already correctly read 'Version 1.1', matching the cover — unlike NCIE-004/005's amendments, this document did not carry the recurring Edition-lag defect). The running header read 'VERSION 1.1' on all 59 rendered pages after every substantive amendment had been applied.	Cover and Governance table corrected to cite appropriate NCIE-001 v1.3 / NCIE-002 v0.5 / NCIE-003 v0.4 / NCIE-004 v1.2 / NCIE-005 v1.3 / NCIE-007 v1.1. Orchestration authority (NCIE-007 v1.1) and Prior baseline (v1.1) rows inserted. New §0.6 Document Amendment History added. Running header corrected to 'VERSION 1.2' and verified across all 59 rendered pages before delivery, checked proactively per the pattern established across the three prior amendments in this series.

Chapters Reviewed, No Substantive Change Required

Chapters 6, 7, 8, 9, 10, 12, 13, 16, 19, 20, 21, 23, 24 were reviewed against the current approved authority chain per the directive's chapter-cluster instructions. Chapters 6-10 are covered by Ch.5's consolidated §5.3 statement rather than individual per-chapter edits (the directive's own content for this cluster is a single repeated principle applied identically across all six chapters). Chapters 12-13 are covered by Ch.11's consolidated §11.3 statement for the same reason. No direct contradiction with the current approved architecture was found in any reviewed chapter, and no genuinely new behavioral consequence was material enough to warrant a separate chapter-level amendment.

Directive Chapter-Mapping Correction

The supplementary directive (items 62-333) labeled several content clusters with chapter numbers that do not match this document's actual chapter titles (e.g. its 'Chapter 16 — Retrieval' cluster, 'Chapter 19 — Evidence integration', 'Chapter 20 — dependency and invalidation', 'Chapter 21 — retention and lifecycle', 'Chapter 22 — privacy and classification', 'Chapter 26 — Context Security & Isolation', 'Chapter 27 — Voice Context', 'Chapter 28 — Privacy/protected identity', 'Chapter 31 — Observability and audit' do not correspond to this document's real chapter 16 [Private/Shared/Institutional Context Boundaries], 19 [Knowledge Versioning], 20 [Knowledge Challenge/Contradiction], 21 [Knowledge Graph Context], 22 [Dependency Graph], 26 [Sensitive Memory/Protected Identity/Voice Context], 27 [Memory Provenance/Attribution/Audit], 28 [Retention/Legal Hold], or 31 [Scalability/Performance/Observability]). Each content cluster was placed in the real chapter whose actual title is the topical match, per the mapping recorded against each AM entry above, rather than the directive's own (evidently mismatched) chapter-number label. This is recorded as a DIRECT SURGICAL CONSISTENCY CORRECTION to the directive's own chapter references, not a change to NCIE-006's locked 36-chapter structure.

Human Review Register Reconciliation

16 (v1.1) + 4 added – 0 removed = 20 (v1.2). Category deltas: Institutional Confirmation Required 6→9 (+3: HR6-4-2, HR6-18-2, HR6-22-1); Security/Sovereignty Confirmation Required 1→2 (+1: HR6-33-2); all other categories unchanged. Blocking 13→17 (+4, all new items are Blocking); Non-Blocking unchanged at 3.

Additional Consistency Amendments Found (Full-Document Contradiction Scan)

A full-document scan for stale current-authority version references (NCIE-001 v1.2, NCIE-002 v0.4, NCIE-003 v0.3, NCIE-004 v1.1, NCIE-005 v1.2 cited as current; NCIE-006 v1.1 cited as current edition) found one occurrence requiring correction: the running header's 'VERSION 1.1' on all 59 pages (AM-DOC). The Document Governance table's Edition field already correctly read 'Version 1.1' before this amendment (unlike NCIE-004/005, which both had a lagging Edition field) — no fix was needed there. No other pre-existing stale or contradicting statement was found. A semantic scan for statements implying Agent Context = Agent Authority, Ephemeral Agent = Ephemeral Provenance loss, Validated AI Result = Institutional Knowledge, Memory = Current Truth, or Recovery = Reauthorization found none; every governing invariant from the directive's §4-5 lists is stated explicitly and preserved throughout. No additional surgical amendment beyond AM-1 through AM-23 and AM-DOC was required.